

Homework 4: simple Airflow pipeline

Installation and initial configuration for Docker is explained in HW3

https://github.com/chap4ev/hse_dc_2024/tree/main/hw3

Let's build a simple data processing pipeline:

1. download dataset from huggingface and save it locally as csv.
I choose titanic dataset as a classic ML dataset
2. read csv and calculate some statistics
3. write data to local file

Configuring Docker

We will need some additional modules:

datasets - to download dataset from huggingface

pandas - to process this dataset

To add these modules we need to build our own docker image.

Add ./Dockerfile:

```
FROM apache/airflow:2.3.0
COPY requirements.txt /
RUN pip install --no-cache-dir "apache-airflow==${AIRFLOW_VERSION}" -r
/requirements.txt
```

And ./requirements.txt :

```
datasets
pandas
```

Now change docker-compose.yaml so it will not use predefined image and build its own:

```
# image: ${AIRFLOW_IMAGE_NAME:-apache/airflow:2.3.0}
build: .
```

Let's build our image

`docker-compose build`

Additionally we will need a separate volume to store our data:

adding data volume to docker-compose.yaml

```
volumes:
  - ./dags:/opt/airflow/dags
  - ./logs:/opt/airflow/logs
  - ./plugins:/opt/airflow/plugins
  - ./data:/data
```

Initialize airflow:

`docker-compose up airflow-init`

And run airflow services:

`docker-compose up`

Writing DAG

DAG definition

The code is placed inside `./dags/` folder as a separate `.py` file

```
from airflow import DAG
from airflow.operators.python_operator import PythonOperator
from airflow.utils.dates import days_ago

from datetime import datetime, timedelta
import os

from datasets import load_dataset
import pandas as pd

default_args = {
    'owner': 'airflow',
    'start_date': days_ago(0),
    'depends_on_past': False,
}

def download_titanic_dataset(path_to_file):
    dataset = load_dataset('lewtun/titanic', split='train')
    df = pd.DataFrame(dataset)
    df.to_csv(path_to_file, index=False)

def calculate_average_age(path_to_file):
    df = pd.read_csv(path_to_file)
    average_age = df['Age'].mean()
    return average_age

def calculate_survival_chance(path_to_file):
    df = pd.read_csv(path_to_file)
    survival_rate = df['Survived'].mean()
    return survival_rate

def save_results(average_age, survival_chance, path_to_file):
    with open(path_to_file, 'w') as f:
        f.write(
            f'Average Age: {average_age}\n'
            f'Survival chance: {survival_chance}'
        )

with DAG(
```

```

'titanic_data_analysis_with_datasets',
description='Download Titanic dataset, calculate average age and survival
chance, and save results',
schedule_interval="@once",
catchup=False,
default_args=default_args,
) as dag:

download_dataset = PythonOperator(
    task_id='download_titanic_dataset',
    python_callable=download_titanic_dataset,
    op_args=['/data/titanic.csv'],
)

calculate_age = PythonOperator(
    task_id='calculate_average_age',
    python_callable=calculate_average_age,
    op_args=['/data/titanic.csv'],
)

calculate_survival = PythonOperator(
    task_id='calculate_survival_chance',
    python_callable=calculate_survival_chance,
    op_args=['/data/titanic.csv'],
)

save_age_results = PythonOperator(
    task_id='save_results',
    python_callable=save_results,
    op_args=[calculate_age.output, calculate_survival.output,
'/data/results.txt'],
)

download_dataset >> [calculate_age, calculate_survival] >> save_age_results

```

This DAG performs a sequence of 4 PythonOperator tasks:

1. download_titanic_dataset:

This task downloads titanic dataset from huggingface and saves it to specified local file

2. calculate_age:

This task reads a local file with a specified dataset and calculates average passengers age.

3. calculate_survival:

This task reads a local file with a specified dataset and calculates the chance of survival for passengers.

4. save_age_results:

This task takes as input the results of calculations of the previous two tasks and writes them to a local file on disk.

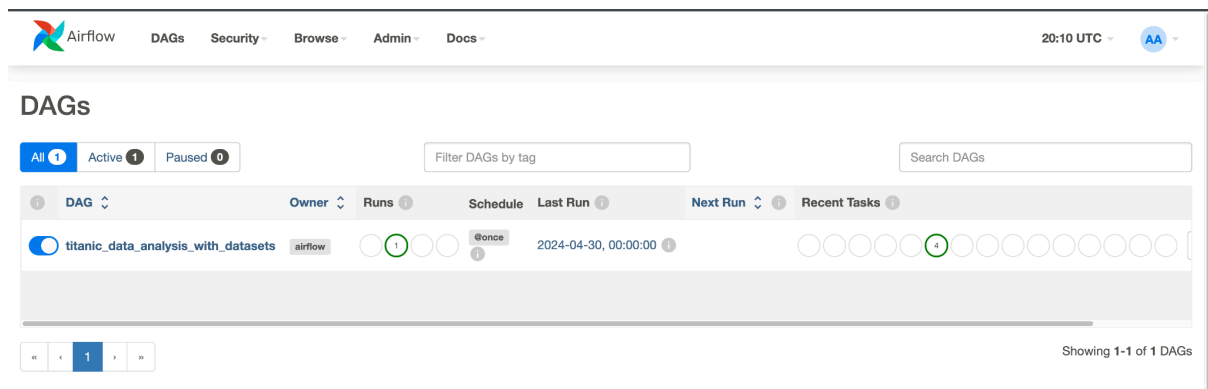
The two counting tasks are independent of each other, so they run in parallel

```
download_dataset >> [calculate_age, calculate_survival] >> save_age_results
```

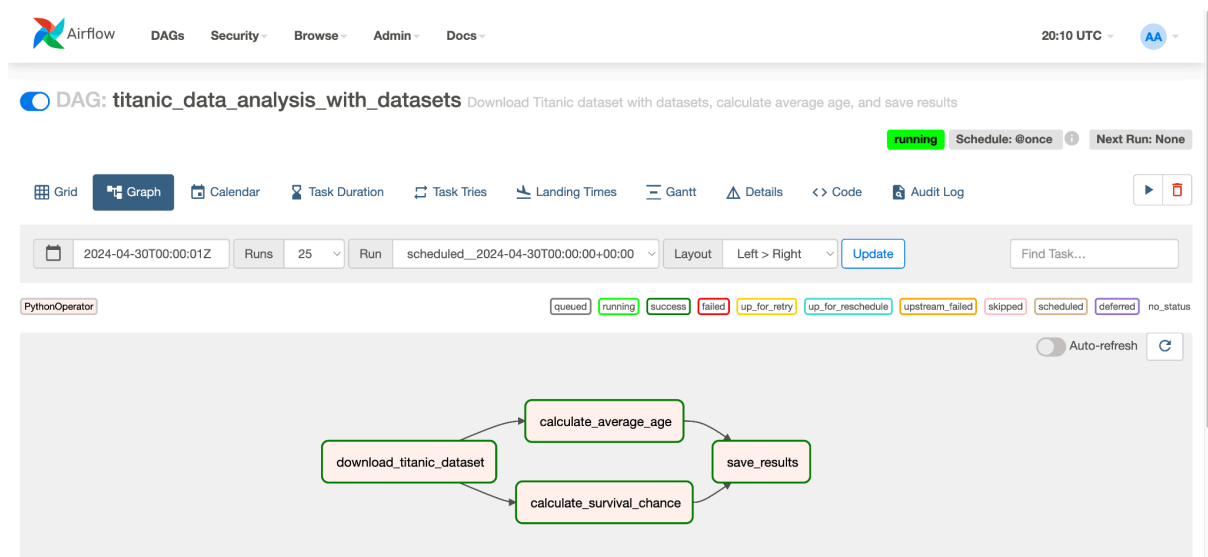
The DAG triggers by hand and does not contain any schedule

Execute DAG

We can see our DAG ready to be executed.



Triggering DAG via “Trigger DAG” button from actions list



And our DAG executed successfully!

We can see tasks execution results as XCOM outputs in interface

Task Instance: calculate_average_age at 2024-04-30, 00:00:00

XCom


Key	Value
return_value	29.69911764705882


Task Instance: calculate_survival_chance at 2024-04-30, 00:00:00


XCom


Key	Value
return_value	0.3838383838383838


The results are stored in our data folder and contain the calculated metrics:


>  dags


>  data


 results.txt


 titanic.csv



 docker-compose.yaml

 Dockerfile

>  logs

>  plugins

 requirements.txt

  results.txt

Average Age: 29.69911764705882
Survival chance: 0.3838383838383838